




Concurrent and Distributed Systems




Running in Parallel - Concurrency

An Example



Process Interactions

- A general classification
 - One to one
 - Any to one
 - One to (one of) many
 - Many to Many



A Classification

- One to one
 - Appropriate in systems with static configurations of interactions between individual processes
 - Example: Pipeline in Unix commands
- Any to one
 - Multiple clients interact with a single server
 - Clients invoke a well known server
 - Server accepts requests from any client
 - Server does not know which client will interact next, waits for the next client
 - Mail server + client, Web server + client



A Classification

- One to many
 - Used to notify a set of interested clients
 - Broadcast (sent out to everyone)
 - Usually no record of reception of communication
 - Clients 'listen' for information
 - Multicast (sent out to a specific set of recipients)
 - How to identify the recipients (clients join a list – mailing list)
 - Reliable, Unreliable (like broadcast)
 - Used in fault tolerant systems




A Classification

- Any to (one of) many
 - Service offered by anonymous server processes
 - Clients requests service from any available server
 - This style usually reduces to one of the other styles
- Many to many
 - Usually implemented by shared data
 - Any number of processes can interact
 - Requires synchronisation to prevent chaos




Forms of Process Interactions

- Co-operation
(shared memory)



- Communication
(message passing)



Implementing IPC

- Shared Memory
 - Processes/threads involved share a **common buffer pool**
 - Buffer can be explicitly implemented by programmer
- Inter-Process Communication without shared memory
 - IPC has at least two operations
 - **Send** (message)
 - **Receive**(message)
 - Messages can be either fixed or variable size
 - A **link** between the involved processes must exist



Implementation of IPC

- Physical implementation
 - Shared memory
 - Hardware bus
 - Network
- Logical implementation of **link** and **send()** & **receive()**:
 - Direct or indirect communication
(naming; processes need to have a way to identify each other)
 - Synchronisation: blocking or non-blocking send/receive
 - Automatic or explicit buffering
 - Send by copy or send by reference
 - Fixed-sized or variable-sized messages
 - continued ... !

Direct communication

- Processes need to explicitly **name** the receptionist / sender (synchronous addressing)
 - Send(P, message)
 - Receive(Q, message)
- Link is established automatically between the two parties
- Processes only need to know each other
- A link is established between exactly two processes
- Between each pair of processes there exists exactly one link
- Disadvantage: limited modularity (changing code)
- Also asynchronous addressing possible
 - Send(P, message) – send a message to process P
 - Receive(id, message) – receive a message from any process; id holds the name of the processes with which communication took place



Indirect Communication

- Messages are send to **mailboxes** or **ports**
- Mailbox is an abstract concept
 - Object into which messages can be included and removed
 - Each mailbox has its unique identification
- Processes can communicate with other processes via different mailboxes
- Communicating processes need to have shared mailboxes
 - Send(A, message) – send a message to mailbox A
 - Receive(A, message) – receive a message from mailbox A
- A link is only established if the processes share a mailbox



Indirect Communication


- A link may be established between more than two processes
- Between a pair of processes there may be any number of links represented by different mailboxes
- How are messages linked to processes?
 - Allow only links between two processes
 - Allow at most one process at a time to execute receive()
 - Allow the system to select which process will receive the message; the system may identify the receiver to the sender

Indirect Communication

- Mailboxes may be owned by
 - a user process
 - Owner process may only receive messages
 - Other processes (users) may only send messages
 - When the owner dies, the mailbox disappears, too.
 - Users need to be notified of the disappearance of a mailbox
 - the Operating System
 - Independent, not associated with any process
 - Operating system offers mechanisms for
 - Creating a new mailbox
 - Send and receive messages
 - Delete a mailbox



Synchronisation




- Message passing may be blocking or non-blocking (synchronous and asynchronous)
- Blocking send
 - Sending process is blocked until the message has been received by the receiving process or mailbox
- Non-blocking send
 - Sending process resumes operation immediately after sending the message
- Blocking receive
 - The receiving process blocks until a message has been received
- Non-blocking receive
 - The receiver retrieves a valid message or a NULL message

Buffering



- Messages exchanged always reside in a temporary queue
- Zero capacity
 - Maximum length 0 \rightarrow no messages can 'wait' in the queue
 - Sender must block until the receiver gets the message
 - Also called a message passing system without buffering
- Bounded capacity
 - Finite length $n \rightarrow$ the queue can hold at most n messages
 - Queue not full: message is stored in the queue (either a copy or a ref); sender can continue execution without waiting
 - Queue full: sender blocks until space is available
- Unbounded capacity
 - Potentially infinite length
 - Sender never blocks

Example – Message Queue



```
import java.util.*;

public class MessageQueue{
    private Vector queue;

    public MessageQueue() {
        queue = new Vector();
    }

    public void send(Object item){
        queue.addElement(item);
    }
}
```

```
public Object receive() {
    Object item;
    if (queue.size() == 0)
        return null;
    else {
        item = queue.firstElement();
        queue.removeElementAt(0);
        return item;
    }
}

private Vector queue;
}
```



Example – Message Queue

- Message Queue for Producer Consumer Example from lecture 3
- Buffer is unbounded and provided by Vector class
- `send()` and `receive()` are non-blocking
- Consumer needs to evaluate the result from `receive()`!
 - message may be NULL



When Shared-Memory is useful

- In an unprotected system where all processes and OS run in the same address space
 - Mac-OS (until 7.5)
- The language operates a simple OS
 - Ada, ML
- In systems where multithreading is provided above the OS
 - Sun LWT library



When Shared Memory is not useful

- In protected systems where processes run in separate address spaces
 - Protection and addressing are orthogonal
- Between processes on different CPUs or machines
 - However, distributed shared memory
- In systems where high flexibility is required
 - Distribute process on different machines not possible!
- In systems where process migration is desirable
 - Migration and shared memory are incompatible
 - However, distributed shared memory



Example

- Play with the alternative producer consumer!?

Summary

- Process communications (1:1, 1:m, ...)
- Shared memory – direct communication
- Synchronisation (blocking/non-blocking)
- Buffers (0, finite, ...)

