

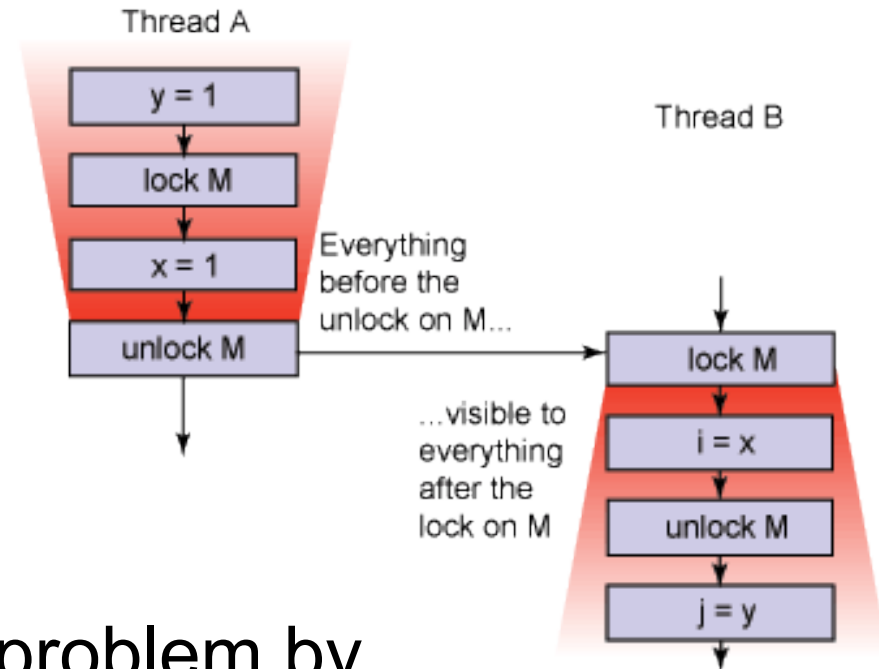


Concurrent and Distributed Systems

Java Synchronisation

Contents

- Background
- Bounded Buffer problem
- Critical Sections
- Solving the critical section problem by synchronising processes / threads
- Java Synchronisation



Background

- Cooperating sequential threads/processes run asynchronously and share data
- Concurrent access to shared data may result in data inconsistency.
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes.
- Illustrate the problem with the Bounded Buffer problem
- Shared-memory solution to bounded-buffer problem has a race condition on the class data ***count***

Bounded Buffer Problem

- Producer

```
while(count == BUFFER_SIZE)
    ; // no-op
// add an item to the buffer
++count;
buffer[in] = item;
in = (in + 1) % BUFFER_SIZE;
```

- Consumer

```
while (count == 0)
    ; // no-op
// remove an item from the buffer
--count;
item = buffer[out];
out = (out + 1) % BUFFER_SIZE
```



Bounded Buffer Problem

- Both consumer and producer work well separately, however, they may not function in combination
- Variable count is shared
- Assume ++count and --count happen concurrently!
- Is the result 4, 5, or 6 (due to processor operations) ?
- (how? Hint: which starting value is possible?)
- Race Condition!

```
register1 = count;  
register1 = register1 + 1;  
count = register1;
```

```
register2 = count;  
register2 = register2 - 1;  
count = register2;
```



Race Conditions: The problem

- Shared variables are written to/read from
- Transfer from one consistent state to another takes several separate operations
- Context switch can happen any time, and operations be interrupted
- Concurrent threads (multi-processor) may share data, also.
- This leads to corrupted data
- Approach: “Critical Section”

Critical Section

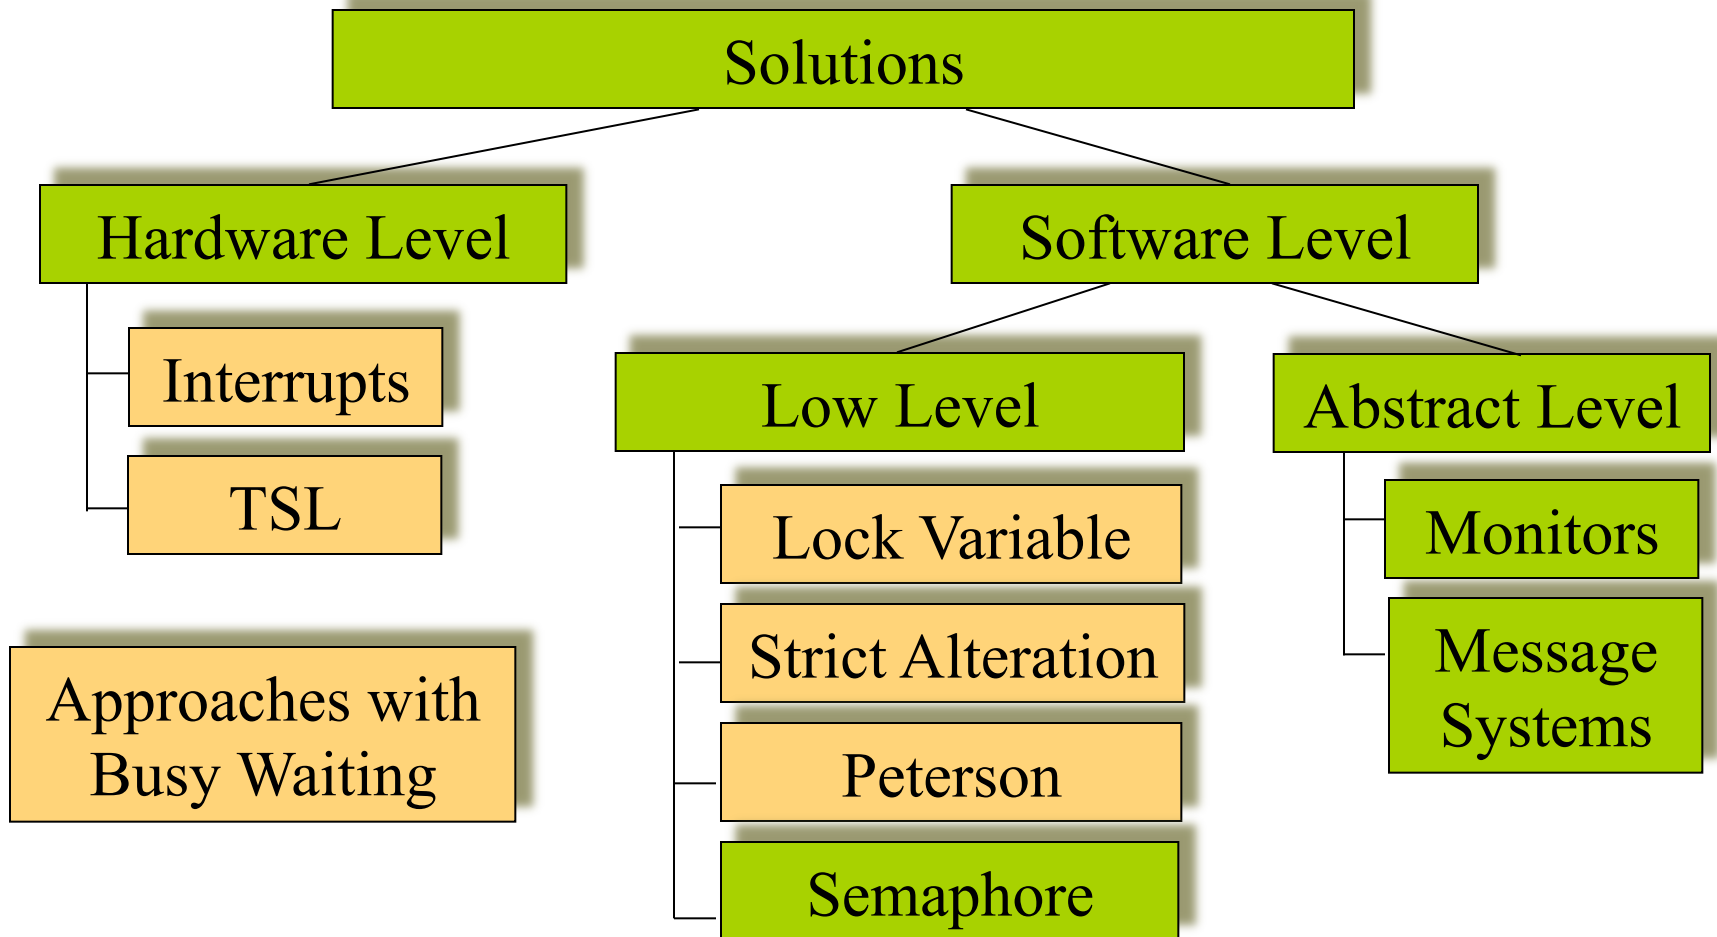
- To prevent race conditions → **only one thread at the time** manipulates the variable count
- Synchronisation of threads is required!
- Sections of code are declared critical
 - Changing common variables
 - Updating a shared table
 - Writing to a file
- Access to critical sections is regulated
 - If one thread executes in a **critical section** no other thread may enter their critical sections
 - Mutual exclusion in time

Solving the Critical Section Problem

- A solution must satisfy four requirements
 - No two processes may be **simultaneously** inside their critical sections
 - No assumptions may be made concerning **speeds or numbers of processors**
 - No process running outside its critical **region may block** other processes
 - No process should have **to wait forever** to enter its critical region (fairness/starvation)



Possible Solutions



Java Synchronisation

- Enforcing mutual exclusion between threads → Thread Safe
- Alternative to Busy Waiting
- Solving Race Conditions
 - Synchronized
 - Wait(), notify(), notifyAll()

Busy Waiting

- Remember Bounded Buffer problem
- → wait until Buffer is not empty / not full
- Alternative: Thread.yield()
 - Thread stays in runnable state
 - Allows JVM to select another thread for execution (equal priority), if any!
 - Problem: Potential **deadlock!!!**



Deadlock scenario

- Deadlock (informal): a series of processes/thread is waiting on conditions (resources) depending on the other processes/thread in the set and no one can run. Permanent condition.

Necessary conditions for deadlock:

1. Mutual exclusion
2. Hold and wait
3. No preemption
4. Circular wait

Deadlock scenario

- JVM uses priorities, thread with highest priority of all threads in runnable state is run before threads with lower priority

- `myThread.setPriority(8);`

almost ever a good solution !

- Producer has higher priority than consumer
 - If buffer is full, producer will execute `yield()`
 - Consumer still cannot run because of lower priority
 - Deadlock!



Fixing Race Conditions

- Java introduces keyword **synchronized**
- Every java *Object* has an associated lock
- Object associated with Bounded Buffer class also has a lock associated
- Normally, when a method is invoked, the lock is ignored
- However, using **synchronized** requires owning the **lock for the object**




Synchronized mechanism

- If the lock is not available (owned by another thread) the thread blocks
- The blocked thread is put into a queue called *entry set*
- Entry set represents all threads waiting for the lock to become available
- The lock is *released when* the owning thread exists a synchronized method
- One thread from the entry set gets the lock



Entry Set



Code example

```
public synchronized void  
    enter(Object item) {  
  
    while (count == BUFFER_SIZE)  
        Thread.yield();  
    ++count;  
    buffer[in] = item;  
    in = (in + 1) % BUFFER_SIZE;  
}
```

```
public synchronized Object remove() {  
    Object item;  
  
    while (count == 0)  
        Thread.yield();  
    --count;  
    item = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    return item;  
}
```

- **Still danger of Deadlock!**

Wait() and Notify()

- Every lock is also equipped with a *wait set*
- If a thread determines it cannot proceed inside a synchronized method it calls wait()
 - Thread releases the lock for the object
 - The state of the thread is set to blocked
 - The thread is placed in the wait set
- Other threads may acquire the lock
- Deadlock is prevented!


Notify()

- Normally when a thread exits a **synchronized** method, it just releases the lock (perhaps removing one thread from the *entry set*)
- Notify()
 - Picks an arbitrary thread T from the *wait set* and puts it into the *entry set*
 - Moves the state of the thread from blocked to runnable
 - T now competes for the lock with all threads in the entry set
 - Once it owns the lock, the wait() call returns

Notify() cont.

- Wait() and Notify() are a synchronisation but even more a communication mechanism
- Wait() and Notify() are independent of the conditions they are used for!
- Wait() and Notify() need to be called **from within** a synchronized block – otherwise **race condition!**
 - `IllegalMonitorState` exception

Entry and wait sets



Code example

```
public synchronized
    void enter(Object item) {
    while (count == BUFFER_SIZE)
        try {
            wait();
        }
        catch (InterruptedException e) { }
    }
    ++count;
    buffer[in] = item;
    in = (in + 1) % BUFFER_SIZE;
    notify();
}
```

```
public synchronized Object remove()
{ Object item;
  while (count == 0)
    try {
        wait();
    }
    catch (InterruptedException e)
    { }
    --count;
    item = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    notify();
    return item;
}
```

Watch the while loop around wait()!

Multiple Notifications

- `notify()` **selects an arbitrary thread from the wait set**. This may not be the thread that you want to be selected.
- Java does not allow you to specify the thread to be selected.
- `notifyAll()` **removes ALL** threads from the wait set and places them in the entry set. This allows the threads to decide among themselves who should proceed next.
- Useful if threads may wait for several conditions
- However, a thread may be woken up for an entirely different condition! - Put `wait()` into a while loop!
- `notifyAll()` is a **conservative** strategy that works best when multiple threads may be in the wait set
- **Inefficient**, since all threads need to re-acquire the lock!



Block Synchronisation

- Blocks of code – rather than entire methods – may be declared as synchronized.
- This yields a lock scope that is typically smaller than a synchronized method.
- Uses a java object to perform the synchronisation
- Used for larger methods where only a small part is a critical section
- Use of wait() and notify() possible (use the same object)
- Useful for static methods



Code example

```
Object mutexLock = new Object();  
...  
public void someMethod() {  
    // non-critical section  
    synchronized(mutexLock) {  
        // critical section  
    }  
    // non-critical section  
}
```

```
Object mutexLock = new Object();  
...  
public void someMethod() {  
    // non-critical section  
    synchronized(mutexLock) {  
        try{  
            mutexLock.wait();  
        }catch (InterruptedException ie() {}  
    }  
    // non-critical section  
}  
public void someOtherMethod() {  
    synchronized(mutexLock) {  
        mutexLock.notify();}  
}
```



Some rules on synchronisation

- A threads that owns the lock for an object may enter another **synchronized** method of the same object
- A thread can nest synchronized method invocations for different objects. Thus a thread can own the lock for several objects
- If a method is not declared synchronized it can be invoked regardless of lock ownership, even while a synchronized method of the same object is being executed
- If the wait set for an object is empty, then a call to notify() or notifyAll() has no effect
- If an exception occurs while a thread holds a lock, the lock is freed → possible inconsistent state in the object