



Concurrent and Distributed Systems

Inter-Process Communication & Synchronisation

Process (thread) Interactions

- Processes may need to *cooperate* to carry out a task
 - A process may make a request and *wait* for the service to be done
 - A process may need to send the requested data, or *signalling* that a task has been done.
- Cooperating processes are in charge of implementing interaction policies, which should fulfil desirable properties, such as absence of deadlock, fairness (non-starvation), ...

We need to keep this in mind when programming
concurrent – either processes or threads - applications

- Competing (cooperating) processes need to *wait* to acquire a shared resource and need to be able to *signal* to indicate that they have finished with that resource



Synchronisation for shared data

- A number of processes are simultaneously accessing memory
 - Reading a memory location is atomic
 - Writing a memory location is atomic
- There may be arbitrary interleaving of machine instruction execution → arbitrary interleaving of memory accesses
- Between any two instructions any number of instructions from any other computation may be executed
- Consider again the usual standard example $x := x + 1$
read + op + write



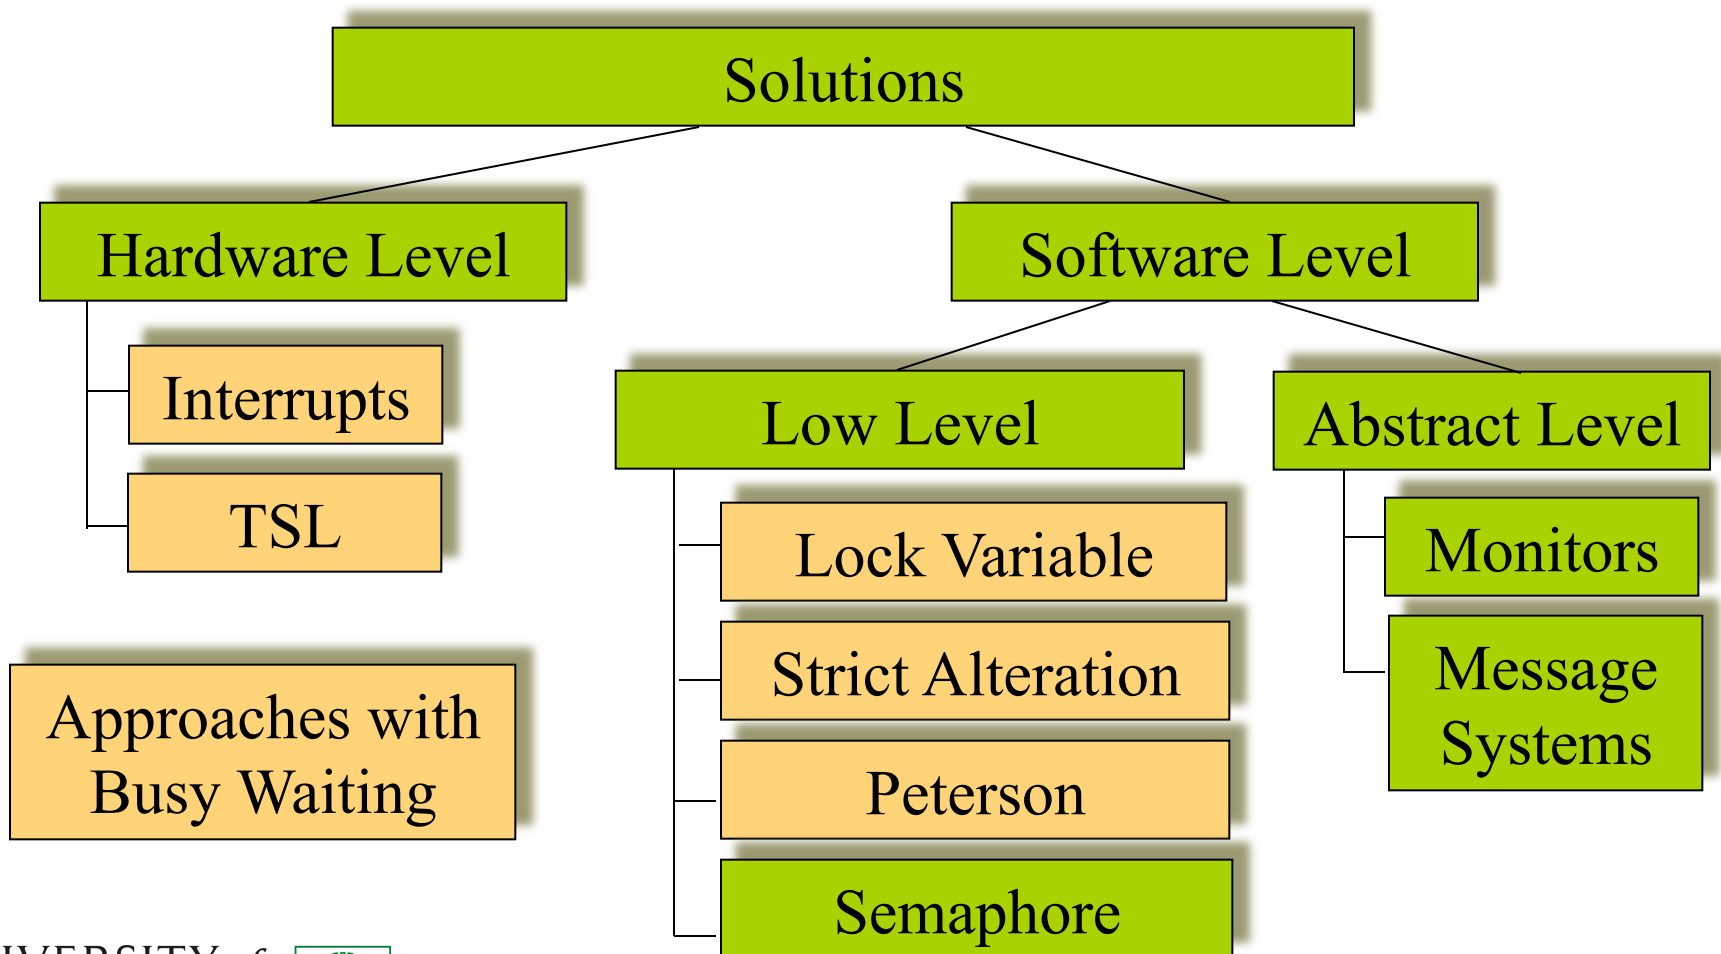
Critical Section Specification

Recall:

1. At **most one** process inside a critical section (mutual exclusion)
2. **No assumptions on speeds** or numbers of processors and processes
 1. Processes may be running on a multiprocessor
 2. Not necessarily a Hardware operation available
3. No process running outside its critical region **may block** other processes
4. No process should have to **wait forever** to enter its critical region
 1. no starvation,
 2. progress – no indefinite busy wait
5. **No deadlock!**



Possible Solutions



1. Disabling Interrupts

Interrupts are the base for timing and preemption and external device handling. Could be disabled:

- No clock interrupts can occur
- No CPU context switch
- No other processes can access the shared resource
- Possible abuse of the system!
- Not applicable to computers with more than one CPU
(disabling interrupts only affects one CPU! processes on other CPU may continue running)
- Not applicable as a general mutual exclusion mechanism!



2. Test and Set Lock (TSL)

- Hardware operation often supported by multiprocessor computers
- Reads the contents of a memory address into a register and writes a nonzero value to that location
- Is register zero? If not some process is in the critical section
- Guaranteed to be indivisible
- No context switch can occur
- Memory bus is blocked to prevent other CPUs from accessing memory during the operation



TSL- implementation example

Uses **tsl** to implement a critical region using flag (memory cell)

enter_region:

| | |
|---------------------------|--|
| tsl register, flag | ;copy flag to register (r) & set flag to 1 (w) |
| cmp register, #0 | ;was flag set? If register = 0 return 0 |
| jnz enter_region | ;if not zero from cmp (flag set) loop |
| ret | ;critical region entered, return to caller |

Any other enter_region call here loops on cmp register, #0

leave_region:

| | |
|--------------|--|
| mov flag, #0 | ;unset flag |
| ret | ;left critical section, return to caller |

(for xi86 CPUs was xchg m, r)



3. Lock Variables (I)

- Single, shared variable; initially set 0
- Enter a critical section: test variable
 - if 0, set it 1, enter critical section
 - If 1, wait until it becomes 0
- Race condition still occurs, if context switch happens after the check (two processes might read 0, then ...)
- No solution (as is)!

Let's see something more elaborated ...

Lock Variables (II)

```
public class ME extends BaseME{
```

```
    public ME() {  
        flag[0] = false;  
        flag[1] = false;  
    }
```

```
    public void enteringCS(int t) {  
        int other = 1 - t;  
        flag[t] = true;  
  
        while (flag[other] == true)  
            Thread.yield();  
    }
```

```
    public void leavingCS(int t) {  
        flag[t] = false;  
    }
```

```
    private volatile boolean[]  
        flag = new boolean[2];  
}
```

- Thread 0 and thread 1
- Possible endless loop:

if context switch occurs
after setting the flag, both
threads wait for the other

4. Strict Alternation

```
public class StrictA extends baseME{  
    public StrictA() {  
        turn = TURN_0;  
    }  
    public void enteringCS(int t) {  
        while (turn != t)  
            Thread.yield();  
    }  
    public void leavingCS(int t) {  
        turn = 1 - t;  
    }  
    private volatile int turn;  
}
```

- `turn` keeps track which thread may enter its critical section
- After the first thread finishes, the second may enter, afterwards the first again
- However, if process one wants to enter and it is process two's turn?
- Solution, if all the processes are equally fast
- Violates the above condition: A thread outside its critical section blocks another thread

5. Peterson's Solution (I)

- Algorithm involves
 - an array, one element per thread, and
 - a flag
- First thread sets its array element and thus indicates interest to enter its CS
- `flag` is set to the other thread
- If the second thread also wants to enter its CS (array element is set), the first thread blocks
- If both threads call `enterCS` simultaneously, one thread overwrites `flag`, → the first thread proceeds and the second enters the CS afterwards
- No endless blocking!



Peterson's Solution (II)

```
public class PSol extends baseME {
```

```
    public PSol() {  
        flag[0] = false;  
        flag[1] = false;  
        turn = TURN_0;  
    }
```

```
    public void enteringCS(int t) {  
        int other = 1 - t;  
        flag[t] = true;  
        turn = other;  
        while ( (flag[other] == true) &&  
                (turn == other) )  
            Thread.yield();  
    }
```

```
    public void leavingCS(int t) {  
        flag[t] = false;  
    }
```

```
    private volatile int turn;  
    private volatile boolean[]  
        flag = new boolean[2];  
}
```

- This solution works!
- based on busy (active) wait: the while runs continuously (if not interrupted), or yield()



Busy Waiting (BW)

- All solutions presented so far used busy waiting to block processes
 - Does not really block the process but it just enters a loop
 - Wastes CPU time
 - Priority inversion problem (in (BW))
 - Two processes L (low priority) and H (high priority)
 - H is run whenever it is in ready state
 - While L is in its critical section, H becomes ready → context switch
 - H wants to enter critical section → busy waiting for L to leave
 - L is never scheduled to run → L never leaves critical section
 - H never progresses!
- BW can be advantageous in multiprocessor systems
 - For short delays
 - No context switch necessary



6. Semaphore

- **Integer variable**, managed by correct implementations of critical sections (in $P()$ and $V()$)
- Construct that **does not need busy waiting**
- Accessed only by two operations
 - **$P()$** – decrement semaphore (Dutch)
 - Process or thread blocks if semaphore will be negative
 - **$V()$** – increment semaphore
- $P()$ and $V()$ are executed **indivisibly** (only one thread or process at the time can modify semaphore by using $P()$ or $V()$)



Using Semaphore

- Counting semaphore (unrestricted values)
 - Can be used to protect a number of resources
 - Semaphore is initialised with the available number
- Binary semaphore (only values of 0 and 1)

Semaphore S; // initialized to 1

P(S);

CriticalSection();

V(S);

- Associate a process queue with a semaphore
- Processes change into waiting state with P() (if negative)
- Processes are woken up from within V() (change state to ready)
- Control is with the CPU scheduler



Java Semaphores

```
public class Semaphore {  
  
    public Semaphore() {  
        value = 0;  
    }  
    public Semaphore(int v) {  
        value = v;  
    }  
  
    public synchronized void P() { /* see next slide */ }  
    public synchronized void V() { /* see next slide */ }  
  
    private int value;  
}
```

P() and V() operations

```
public synchronized void P() {  
    while (value <= 0) {  
        try {  
            wait();  
        }  
        catch (InterruptedException e) { }  
    }  
    value --;  
}
```

```
public synchronized void V() {  
    ++value;  
  
    notify();  
}
```



Example using Semaphore

```
public class SemaphoreExample {  
  
    public static void main(String args[]) {  
        Semaphore sem = new Semaphore(1); // get a semaphore & initialise 1  
        Worker[] bees = new Worker[5];   // get 5 threads  
  
        for (int i = 0; i < 5; i++)  
            bees[i] = new Worker(sem);    // provide semaphore to threads  
  
        for (int i = 0; i < 5; i++)  
            bees[i].start();              // start threads  
  
    } // end main  
} // end class
```



Example using Semaphore

```
public class Worker extends Thread {  
  
    public Worker(Semaphore s) { sem = s;}           // constructor  
    public void run() {  
        while (true) {  
            sem.P();           // enter critical section, may block here  
                               // in critical section  
            sem.V();           // leave cs and wakeup other threads  
                               // out of critical section  
        } // end loop  
    } // end run method  
  
    private Semaphore sem;  
}
```



Homework

Use semaphore(s) to implement a rendezvous.



Deadlock with P() and V()

- **Deadlock** (of a set of processes)
 - two or more processes are waiting indefinitely for an event that can only be caused by one of the waiting processes.
- Let S and Q be two semaphores initialised to 1

| | |
|----------|----------|
| P_0 | P_1 |
| P(S); | P(Q); |
| P(Q); | P(S); |
| \vdots | \vdots |
| V(S); | V(Q); |
| V(Q) | V(S); |



- **Starvation** (of one or more processes)
 - indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended. (LIFO queue)

7. Monitors

- Semaphores are error prone!
 - Correct order is essential
 - Errors are hard to detect, depend on particular execution sequence
 - Swap the order of $P()$ and $V()$
 - Replace $V()$ with $P()$
 - Omit either $P()$ or $V()$
 - Consider major, large-scale software development projects
- Monitors are a high-level construct to prevent such errors



Monitors

- Monitor presents a set of programmer defined operations that provide **automatic mutual exclusion**
- Monitor type also contains variables to define the **state of an instance** of the monitor
- A monitor method can **only access monitor internal data and formal parameters**
- Local variables may only be accessed from within the monitor
- Monitor construct **prohibits concurrent access** to all methods defined within that monitor

Monitors

```
monitor Monitor-name
{
    integer i;    // variables
    condition c; // condition variables
    public producer(...) {
        ...
    }
    public consumer(...) {
        ...
    }
}
```




Condition Variables

- Condition variables are used for user specific synchronisation (buffer full/empty) condition x , y ;
- Operations `wait()` and `signal()` are defined
 - `x.wait()` suspends the invoking thread until `x.signal()` is called by another thread
- Thread frees the monitor after blocking
- After signalling, a thread:
 - Signal-and-wait: signalling thread waits for other thread to finish in the monitor or to block on another condition
 - Signal-and-continue: signalling thread continues processing. The woken up thread continues afterwards.



Monitors



Message Passing (recap)

- Consider distributed systems **without shared memory**
 - Semaphores are too low level
 - Monitors are inapplicable
 - No information exchange possible between machines
- Message Passing!



Message Passing

- Implements two-way messages
 - Send(destination, message)
 - Receive(source, message)
- Implemented as system calls rather than language constructs (no shared memory)
 - Messages are buffered by the operating system
 - Usually provided by a library
- Receiver may block when there are no messages or return with an error code
- Issues
 - Messages may be lost by the network
 - Naming of processes
 - Authentication of processes
 - On the same machine: performance



Producer – Consumer Example

- ! No shared memory – possible implementation
- All messages are the same size
- Messages are buffered by the OS
- N messages are used (N elements in a buffer)
- Consumer starts by sending N empty messages to the producer
- Whenever the producer has an item to send, it takes an empty message, fills it, and sends it to the consumer
- Producer will be blocked if there are no empty messages waiting
- Consumer blocks if there are no filled messages waiting
- Zero buffer option possible



Producer-Consumer Problem

```
#define N 100                                /* number of slots in the buffer */

void producer(void)
{
    int item;
    message m;                                /* message buffer */

    while (TRUE) {
        item = produce_item();                /* generate something to put in buffer */
        receive(consumer, &m);                /* wait for an empty to arrive */
        build_message(&m, item);              /* construct a message to send */
        send(consumer, &m);                   /* send item to consumer */
    }
}

void consumer(void)
{
    int item, i;
    message m;

    for (i = 0; i < N; i++) send(producer, &m); /* send N empties */
    while (TRUE) {
        receive(producer, &m);                /* get message containing item */
        item = extract_item(&m);              /* extract item from message */
        send(producer, &m);                  /* send back empty reply */
        consume_item(item);                  /* do something with the item */
    }
}
```

